Statistical Methods of Machine Learning
Assignment 1

Martin Grunbaum
Martin Jorgensen
Thomas Barnholdt

February 20, 2014

I.1.1.1

Given

$$a = \begin{pmatrix} 1 \\ 2 \\ 2 \end{pmatrix} b = \begin{pmatrix} 3 \\ 2 \\ 1 \end{pmatrix}$$

Then $a^T b = 1*3 + 2*2 + 2*1 = 3 + 4 + 2 = 9$

I.1.1.2

The *l2-norm* or *Euclidean norm* $||a|| = \sqrt{1^2 + 2^2 + 2^2} = 3$

I.1.1.3

The outer product

$$ab^T = \begin{bmatrix} 1*3 & 1*2 & 1*1 \\ 2*3 & 2*2 & 2*1 \\ 2*3 & 2*2 & 2*1 \end{bmatrix} = \begin{bmatrix} 3 & 2 & 1 \\ 6 & 4 & 2 \\ 6 & 4 & 2 \end{bmatrix}$$

I.1.1.4

As M is a diagonal matrix the inverse matrix of M is

$$M^{-1} = \begin{bmatrix} 1/1 & 0 & 0 \\ 0 & 1/4 & 0 \\ 0 & 0 & 1/2 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 0.25 & 0 \\ 0 & 0 & 0.5 \end{bmatrix}$$

I.1.1.5

The matrix-vector product $Ma = \begin{pmatrix} 1*1 + 0*2 + 0*2 \\ 0*1 + 4*2 + 0*2 \\ 0*1 + 0*2 + 2*2 \end{pmatrix} = \begin{pmatrix} 1 \\ 8 \\ 2 \end{pmatrix}$

I.1.1.6

$$A^T = (ab^T)^T = \begin{bmatrix} 3 & 6 & 6 \\ 2 & 4 & 4 \\ 1 & 2 & 2 \end{bmatrix}$$

I.1.1.7

The rank of $A = 1$, because the rows are linearly dependent. We can verify this by observing that the first row can produce the second and third rows with a multiple, e.g. the second row (6 4 2) is the same as the first row (3 2 1) x 2.

I.1.1.8

As A is not full rank, it is not invertible.

I.1.2.1

The derivative of $f(w) = (wx + b)^2$ with respect to w is

$$\begin{aligned} ((wx + b)^2)' &= (w^2x^2 + 2wxb + b^2)' \\ &= 2x^2w + 2xb \\ &= 2x(xw + b) \end{aligned}$$

I.1.2.2

In general

$$\left(\frac{f}{g}\right)'(x) = \frac{f'(x) \cdot g(x) - f(x) \cdot g'(x)}{(g(x))^2}$$

Therefore, differentiating for w we get:

$$f(x) = 1$$
$$f'(x) = 0$$
$$g(x) = (wx + b)^2$$
$$g'(x) = 2x(wx + b)$$
$$\left(\frac{f}{g}\right)'(w) = \frac{0 \cdot (wx + b)^2 - 1 \cdot 2x(wx + b)}{((wx + b)^2)^2}$$
$$= \frac{-1 \cdot 2x(wx + b)}{(wx + b)^4}$$
$$= \frac{-2x}{(wx + b)^3}$$

I.1.2.3

In general

$$(f \cdot g)'(x) = f'(x) \cdot g(x) + f(x) \cdot g'(x)$$

Therefore, differentiating for x we get:

$$f(x) = x$$
$$f(x)' = 1$$
$$g(x) = e^x$$
$$g(x)' = e^x$$
$$(f \cdot g)'(x) = 1e^x + xe^x$$

I.2.1

The plots with gaussian distributions for (μ, σ) pairs $(-1, 1)$, $(0, 2)$ and $(2, 3)$ can be seen in Figure 1. The code for generating the plots can be found in `I_2_1.m`, and the code for our gaussian distribution function can be found in `unigauss.m`.
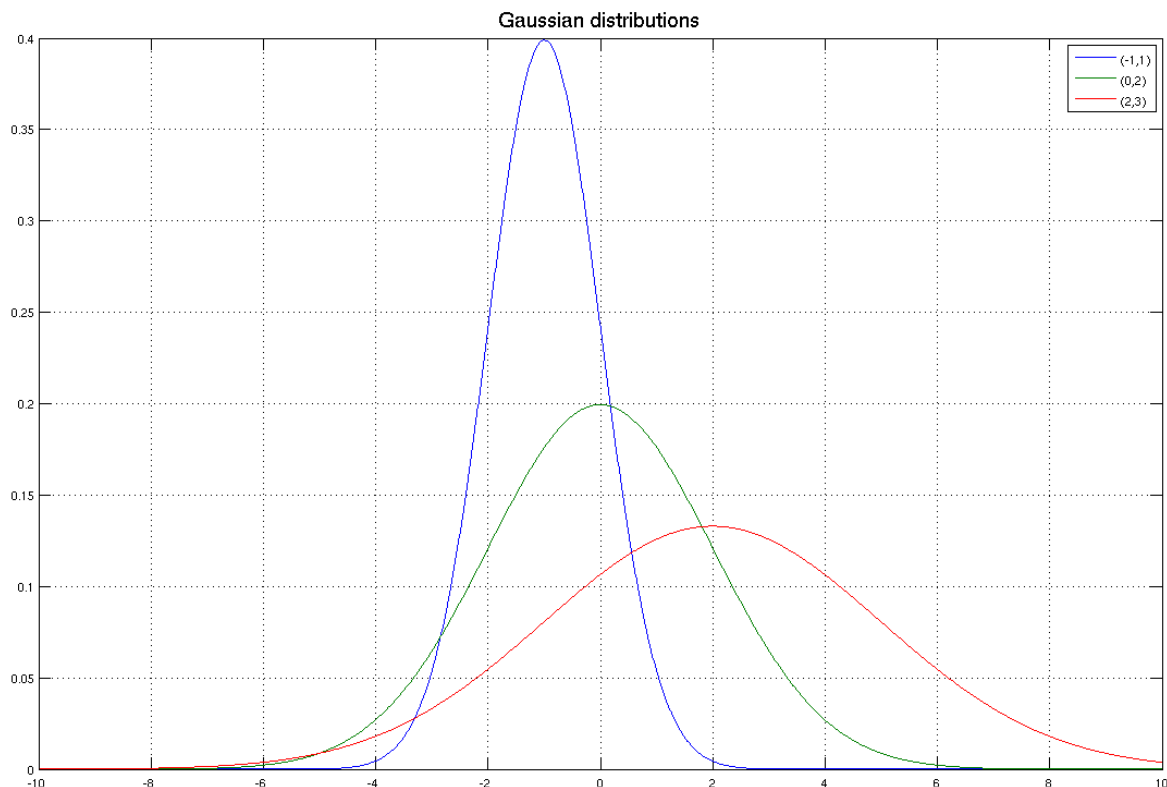

Figure 1: Gaussian distributions plotted with different values for (μ, σ).

I.2.2

Source code is available in `multigauss.m` and `I_2_2.m`. Plot can be seen in Figure 2.

I.2.3

The $l2$ norm of x is

$$mean = \begin{pmatrix} 1 & 2 \end{pmatrix}^T$$

$$\mu = \begin{pmatrix} 1.0006 & 1.9834 \end{pmatrix}^T$$

$$||x|| = l2(mean - \mu) = 0.0366$$

where $l2()$ is a function that calculates the *Euclidean norm* or $l2$ norm of the vector $mean - \mu$.
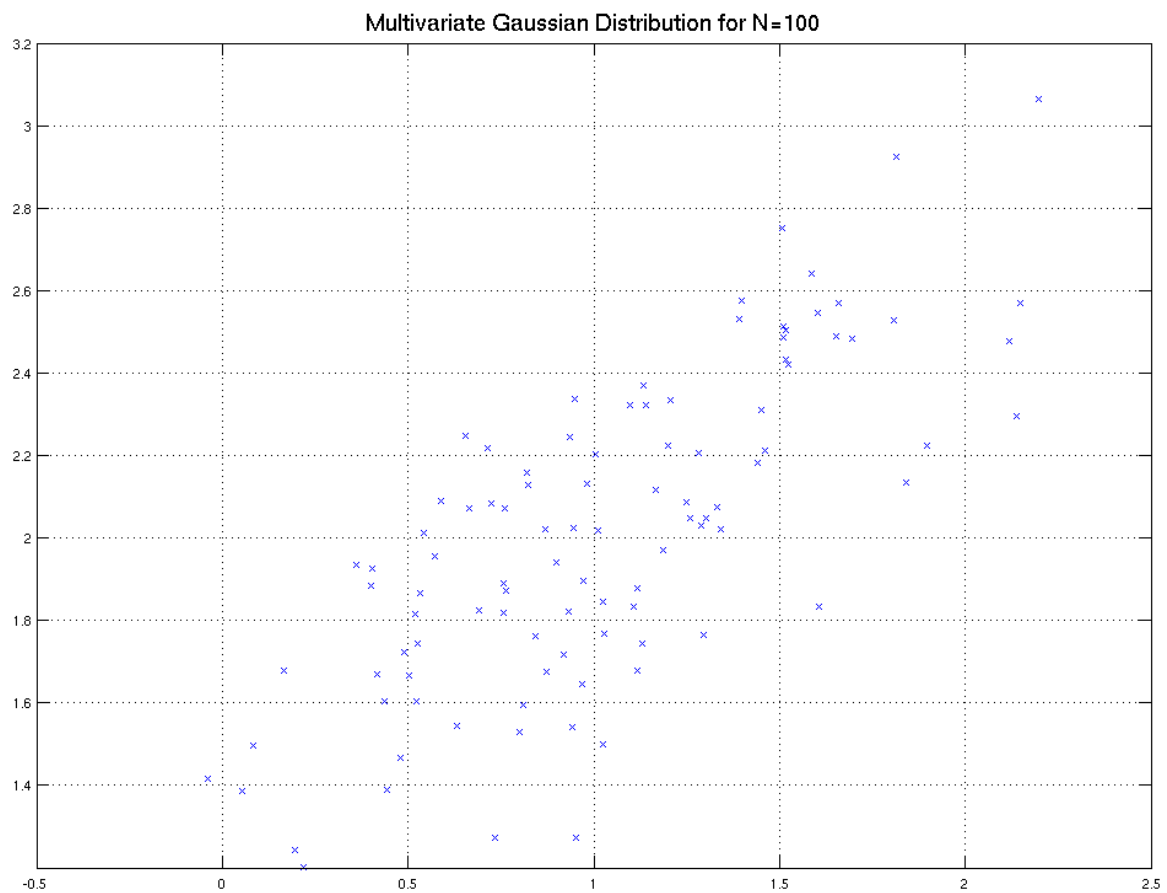
Figure 2: 100 points drawn from a 2-dimensional Multivariate gaussian distribution.

Figure 3 plots the points drawn along with a red circle for the calculated mean and a green circle for μ. There is a difference between the two because the mean is calculated based on the generated data drawn from the multivariate gaussian distribution at random. If we had a number of points approaching infinite, the difference would approach $\bar{0}$. The source code for this excersize can be found in `I_2_3.m`.
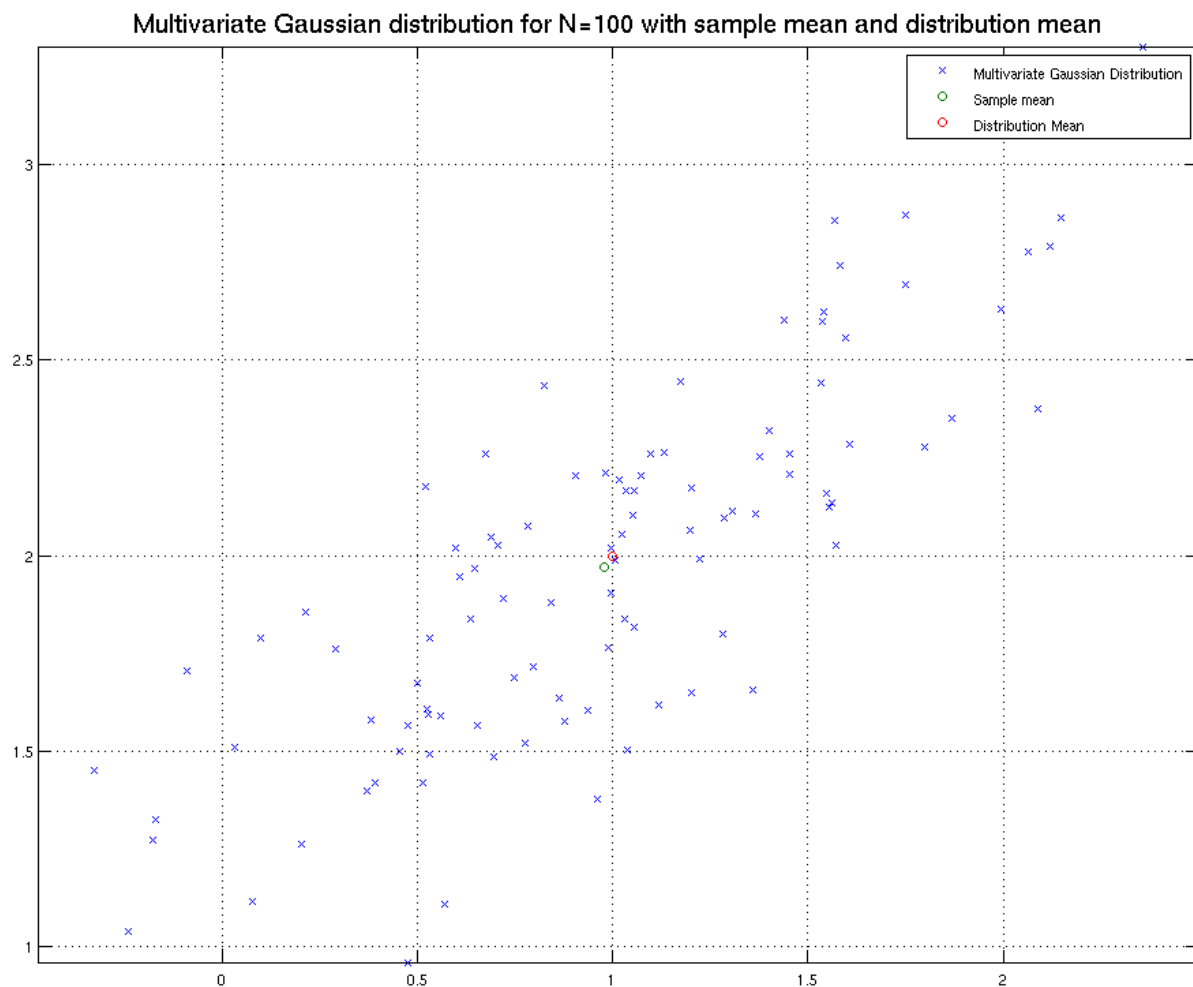
Figure 3: 100 points drawn from a 2-dimensional Multivariate gaussian distribution, plotted with the mean of the points and of μ.

I.2.4

The covariance matrix is full rank 2 and thus has two eigenvectors and eigenvalues. Each eigenvector represents a principal component (or linearly uncorrelated variable), and each eigenvalue a scalar representing the variance. Intuitively, the eigenvectors form a scaled and translated coordinate system centered at the mean of the multivariate Gaussian distribution (μ). If an eigenvalue is 0, the dimensionality is reduced by one. The larger of the two eigenvector/value pairs represents the direction where the ellipsis is widest. The other represents where the ellipsis is narrowest.

The covariance matrix we calculated can be found in Eq 1.

$$\Sigma_{ML} = \frac{1}{N} \sum_{n=1}^{N} (x_n - \mu_{ML})(x_n - \mu_{ML})^T$$
$$= \begin{pmatrix} 0.3239 & 0.2093 \\ 0.2093 & 0.2080 \end{pmatrix} \tag{1}$$

Figure 4 shows a plot of the Multivariate gaussian distribution, plotted with the mean, μ and the two eigenvectors centered in the distribution μ. Figure 5 shows a plot of the 3 rotated distributions along with the distribution rotated to match the largest eigenvector along the x-axis. The angle needed for this was -37.2564^o in our case. Source code is available in `multigauss.m` and `I_2_4.m`. The angle is calculated by by the formula

$$-\arctan(\mathrm{eig}(\Sigma_{\mathrm{ML}}))$$

which return the angle needed for rotating the distribution so its eigen vector is parallel to the x-axis.
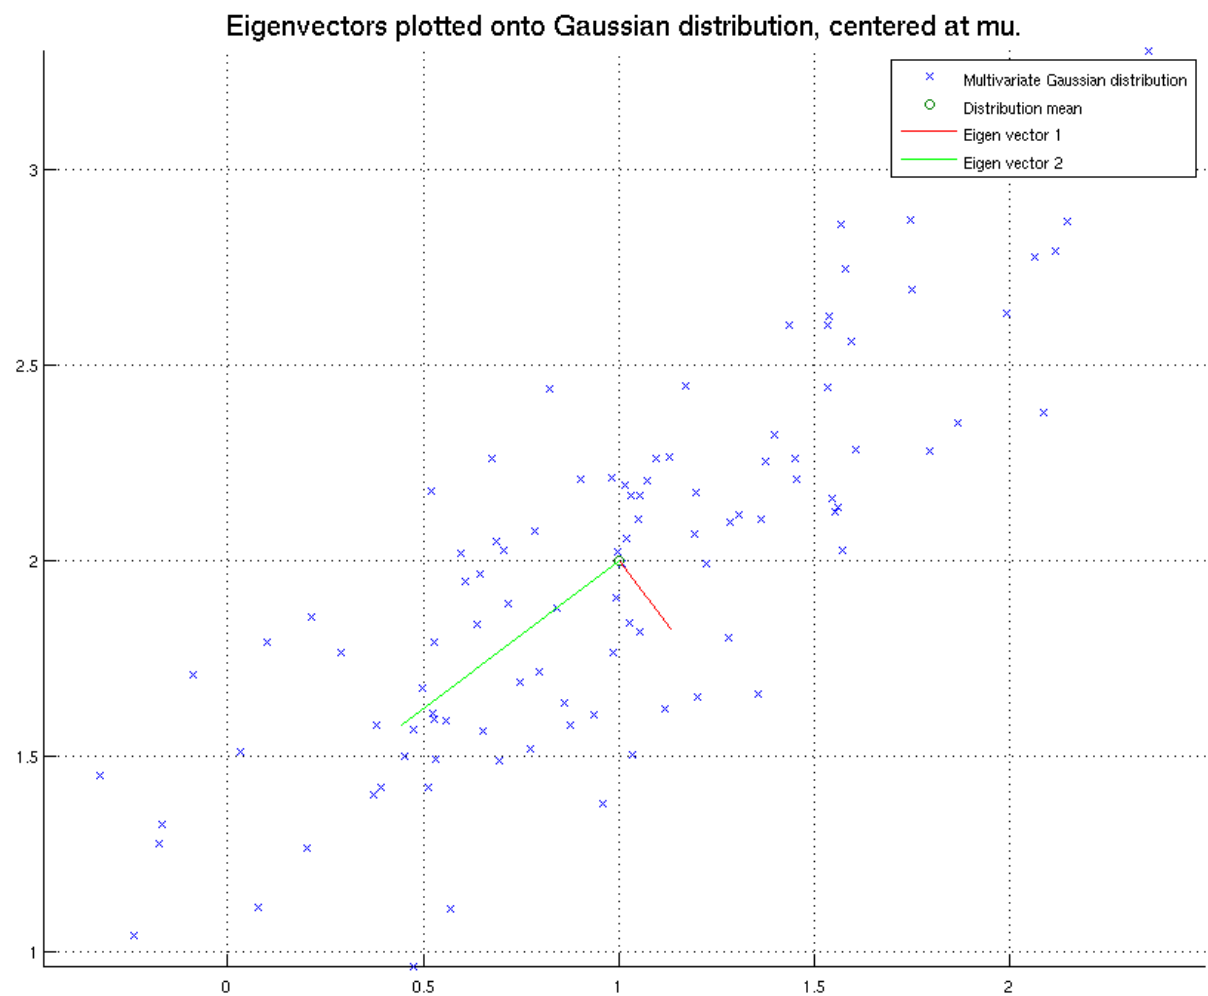
Figure 4: 100 points drawn from a 2-dimensional Multivariate gaussian distribution, plotted with the mean of the distribution, the value of μ and the two eigenvectors centered in the distribution μ.
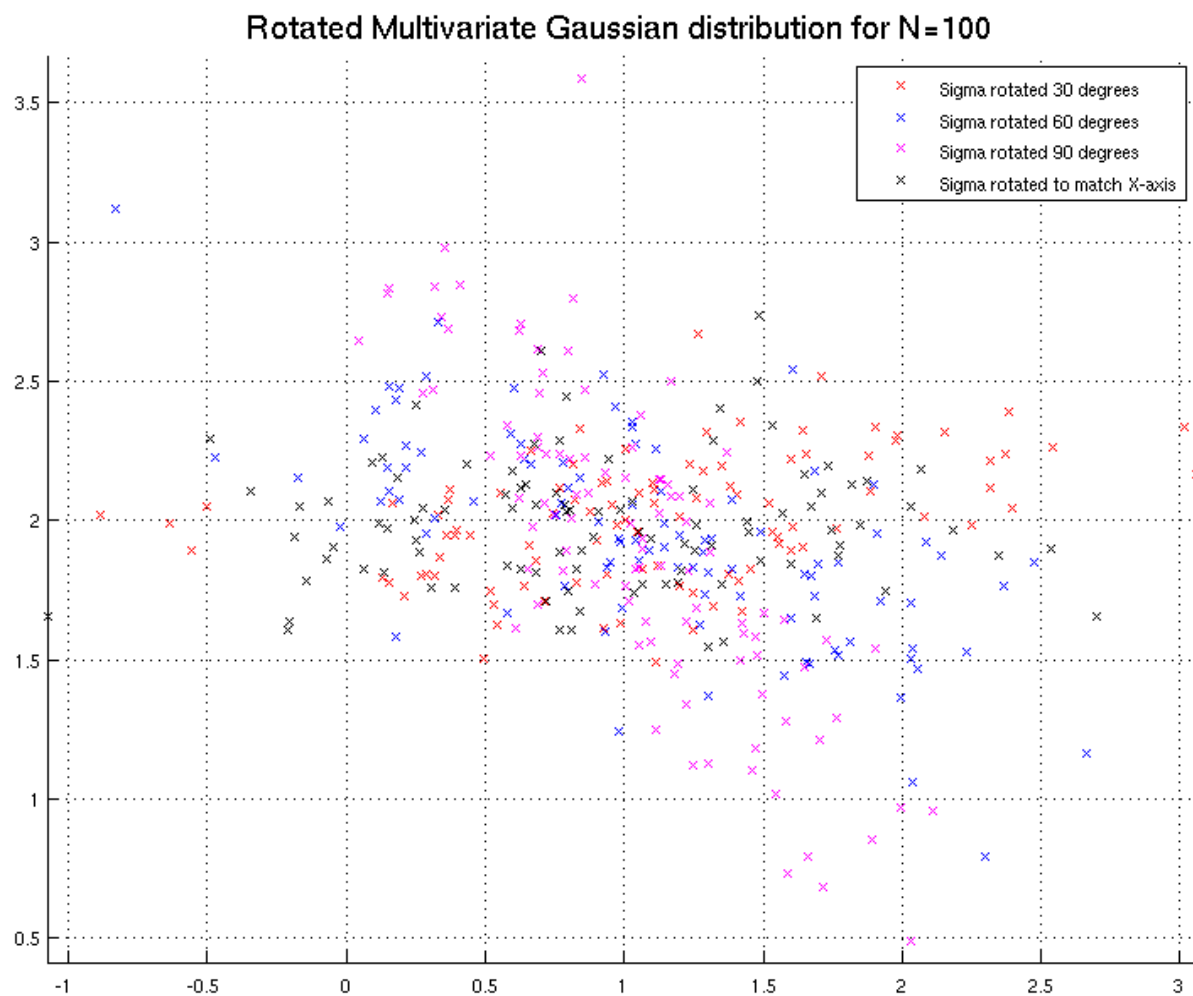
Figure 5: 100 points drawn from a 2-dimensional Multivariate gaussian distribution, rotated at 30, 60 and 90 degrees and lastly also aligned along the x-axis, all distributions in their own color.

I.3

This is the make-up assigment.

Let:

$$\mathcal{N}(x|\mu_A, \sigma_A^2) = \frac{1}{\sqrt{2\pi\sigma_A^2}} e^{-\frac{1}{2\sigma_A^2}(x-\mu_A)^2}$$

$$\mathcal{N}(x|\mu_B, \sigma_B^2) = \frac{1}{\sqrt{2\pi\sigma_B^2}} e^{-\frac{1}{2\sigma_B^2}(x-\mu_B)^2}$$

be univariate Gaussian probability distribution functions.

We look at:

$$f(x) = \mathcal{N}(x|\mu_A, \sigma_A^2) \cdot \mathcal{N}(x|\mu_B, \sigma_B^2)$$
$$= \frac{1}{\sqrt{2\pi\sigma_A^2}} e^{-\frac{1}{2\sigma_A^2}(x-\mu_A)^2} \cdot \frac{1}{\sqrt{2\pi\sigma_B^2}} e^{-\frac{1}{2\sigma_B^2}(x-\mu_B)^2}$$

and want to show $f(x)$ is an unnormalized Gaussian distribution.

As it is to be unnormalized we ignore the constants in front:

$$f(x) = C_1 \cdot e^{-\frac{1}{2\sigma_A^2}(x-\mu_A)^2} \cdot e^{-\frac{1}{2\sigma_B^2}(x-\mu_B)^2}.$$

In the following we alter this expression.

$$e^{-\frac{1}{2\sigma_A^2}(x-\mu_A)^2} \cdot e^{-\frac{1}{2\sigma_B^2}(x-\mu_B)^2} = e^{-\frac{1}{2\sigma_A^2}(x-\mu_A)^2 - \frac{1}{2\sigma_B^2}(x-\mu_B)^2}$$

We only look at the exponent.

$$-\frac{1}{2\sigma_A^2}(x-\mu_A)^2 - \frac{1}{2\sigma_B^2}(x-\mu_B)^2 = -\frac{(x-\mu_A)^2}{2\sigma_A^2} - \frac{(x-\mu_B)^2}{2\sigma_B^2}$$
$$= -\frac{2\sigma_B^2 \cdot (x-\mu_A)^2}{2\sigma_A^2 \cdot 2\sigma_B^2} - \frac{2\sigma_A^2 \cdot (x-\mu_B)^2}{2\sigma_A^2 \cdot 2\sigma_B^2}$$
$$= -\frac{2\sigma_B^2 \cdot (x-\mu_A)^2 + 2\sigma_A^2 \cdot (x-\mu_B)^2}{2\sigma_A^2 \cdot 2\sigma_B^2}$$

We now look at the nummerator.

$$
\begin{aligned}
& 2\sigma_B^2 \cdot (x - \mu_A)^2 + 2\sigma_A^2 \cdot (x - \mu_B)^2 \\
&= 2\sigma_B^2 \cdot x^2 + 2\sigma_B^2 \cdot \mu_A^2 - 2\sigma_B^2 \cdot 2x\mu_A + 2\sigma_A^2 \cdot x^2 + 2\sigma_A^2 \cdot \mu_B^2 - 2\sigma_A^2 \cdot 2x\mu_B \\
&= (2\sigma_B^2 + 2\sigma_A^2) \cdot x^2 + 2\sigma_B^2 \cdot \mu_A^2 + 2\sigma_A^2 \cdot \mu_B^2 - 2 \cdot (2\sigma_B^2\mu_A + 2\sigma_A^2\mu_B) \cdot x \\
&= (2\sigma_B^2 + 2\sigma_A^2) \cdot x^2 + 2\sigma_B^2 \cdot \mu_A^2 + 2\sigma_A^2 \cdot \mu_B^2 - \frac{(2\sigma_B^2 + 2\sigma_A^2)}{(2\sigma_B^2 + 2\sigma_A^2)} \cdot 2 \cdot (2\sigma_B^2\mu_A + 2\sigma_A^2\mu_B) \cdot x \\
&= (2\sigma_B^2 + 2\sigma_A^2) \cdot x^2 + 2\sigma_B^2 \cdot \mu_A^2 + 2\sigma_A^2 \cdot \mu_B^2 - 2(2\sigma_B^2 + 2\sigma_A^2) \cdot \frac{(2\sigma_B^2\mu_A + 2\sigma_A^2\mu_B)}{(2\sigma_B^2 + 2\sigma_A^2)} \cdot x \\
&= 2\sigma_B^2 \cdot \mu_A^2 + 2\sigma_A^2 \cdot \mu_B^2 - \left(\frac{(2\sigma_B^2\mu_A + 2\sigma_A^2\mu_B)}{(2\sigma_B^2 + 2\sigma_A^2)} \right)^2 + (2\sigma_B^2 + 2\sigma_A^2) \cdot \left(x - \frac{(2\sigma_B^2\mu_A + 2\sigma_A^2\mu_B)}{(2\sigma_B^2 + 2\sigma_A^2)} \right)^2 \\
&= C_2 + (2\sigma_B^2 + 2\sigma_A^2) \cdot \left(x - \frac{(2\sigma_B^2\mu_A + 2\sigma_A^2\mu_B)}{(2\sigma_B^2 + 2\sigma_A^2)} \right)^2
\end{aligned}
$$

Where C_2 is independent of x. Setting:

$$
\mu_{AB} = \frac{(2\sigma_B^2\mu_A + 2\sigma_A^2\mu_B)}{(2\sigma_B^2 + 2\sigma_A^2)}
$$

we get:

$$
C_2 + (2\sigma_B^2 + 2\sigma_A^2) \cdot (x - \mu_{AB})^2.
$$

Now we reenter the denominator.

$$
\begin{aligned}
& -\frac{C_2 + (2\sigma_B^2 + 2\sigma_A^2) \cdot (x - \mu_{AB})^2}{2\sigma_A^2 \cdot 2\sigma_B^2} \\
&= -\frac{C_2}{2\sigma_A^2 \cdot 2\sigma_B^2} - \frac{(2\sigma_B^2 + 2\sigma_A^2) \cdot (x - \mu_{AB})^2}{2\sigma_A^2 \cdot 2\sigma_B^2} \\
&= C_3 - \frac{1}{2} \cdot \frac{1}{\frac{(2\sigma_B^2 + 2\sigma_A^2)}{2\sigma_A^2 \cdot \sigma_B^2}} \cdot (x - \mu_{AB})^2
\end{aligned}
$$

Setting:

$$
\sigma_{AB}^2 = \frac{(2\sigma_B^2 + 2\sigma_A^2)}{2\sigma_A^2 \cdot \sigma_B^2}
$$

we get:

$$C_3 - \frac{1}{2} \cdot \frac{1}{\sigma_{AB}^2} \cdot (x - \mu_{AB})^2.$$

We now reenter the base e.

$$e^{C_3 - \frac{1}{2\sigma_{AB}^2} \cdot (x - \mu_{AB})^2} = e^{C_3} \cdot e^{-\frac{1}{2\sigma_{AB}^2} \cdot (x - \mu_{AB})^2}$$
$$= C_4 \cdot e^{-\frac{1}{2\sigma_{AB}^2} \cdot (x - \mu_{AB})^2}$$

And the ignored constant:

$$f(x) = C_1 \cdot C_4 \cdot e^{-\frac{1}{2\sigma_{AB}^2} \cdot (x - \mu_{AB})^2}$$
$$= C_5 \cdot e^{-\frac{1}{2\sigma_{AB}^2} \cdot (x - \mu_{AB})^2}.$$

And we are now on the form we want, and we can se $\mu = \mu_{AB}$ and $\sigma = \sigma_{AB}$.

Old I.3

From here we have the old assignment which haven been altered.

Given:

$$\mu = \begin{pmatrix} \mu_a \\ \mu_b \\ \mu_c \end{pmatrix} \quad x = \begin{pmatrix} x_a \\ x_b \\ x_c \end{pmatrix} \quad \Sigma = \begin{bmatrix} \Sigma_{aa} & \Sigma_{ab} & \Sigma_{ac} \\ \Sigma_{ba} & \Sigma_{bb} & \Sigma_{bc} \\ \Sigma_{ca} & \Sigma_{cb} & \Sigma_{cc} \end{bmatrix}$$

We wish to discover an expression for the conditional distribution $p(x_a|x_b)$ in which x_c has been marginalized out.

We now use, that a vector of length i and a vector of length j can be seen as a vector of length $k = i + j$, and similarly with matrices.

Let:

$$\mu = \begin{pmatrix} \mu_d \\ \mu_c \end{pmatrix} \quad x = \begin{pmatrix} x_d \\ x_c \end{pmatrix} \quad \Sigma = \begin{bmatrix} \Sigma_{dd} & \Sigma_{dc} \\ \Sigma_{dc} & \Sigma_{cc} \end{bmatrix}$$

now we have from chapter 2.3.2 that:

$$\mathbb{E}[x_d] = \Sigma_{dd}$$

$$= \begin{bmatrix} \Sigma_{aa} & \Sigma_{ab} \\ \Sigma_{ba} & \Sigma_{bb} \end{bmatrix}$$

$$\mathrm{cov}[x_d] = \mu_d$$

$$= \begin{pmatrix} \mu_a \\ \mu_b \end{pmatrix}$$

when x_c has been marginalized out.

Now we have from chapter 2.3.1 for the conditional distribution $p(x_a|x_b)$:

$$\mu_{a|b} = \mu_a + \Sigma_{ab}\Sigma_{bb}^{-1}(x_b - \mu_b)\Sigma_{a|b} \qquad\qquad = \Sigma_{aa} - \Sigma_{ab}\Sigma_{bb}^{-1}\Sigma_{ba}.$$

Description	K-value	Accuracy in %
Run on training data	1	100%
Run on test data	1	81.5%
Run on training data	3	86.0%
Run on test data	3	81.5%
Run on training data	5	83.0%
Run on test data	5	68.4%

Table 1: The results from I.4

1   I.4

1.1   I.4.1

The result of our KNN implementation for different k-values and datasets is shown in Table 1. The code to run this particular experiment is in `I_4_1.m`.

With $K = 1$ and running against the training set, the accuracy is 100% since any entry will be matched against itself, and only itself. We also see a general loss of accuracy as K increases. This may be because the point gets matched up against a larger and larger set of the total points, and if there is inherent density clusters in the data then we risk going further and further out as K increases.

1.2   I.4.2

The code for this experiment is in `I_4_2.m`. It uses several auxiliary files found in the same directory. Given a set of possible k values 'PossibleKValues' the following pseudocode runs cross-validation to find the average loss experienced amongst each of the five folds of a cross-validation for each given k value, and then selects the best k as the one with the lowest average loss:

for k in PossibleKValues **do**
  subsets $\leftarrow$ shuffleSplit(dataset,5)
  for $cv = 1$ to 5 **do**
    test $\leftarrow$ subsets[cv]   ▷ bucketJoiner function
    train $\leftarrow$ subsets $-$ test   ▷ bucketJoiner function
    pred $\leftarrow$ kNN(k, train.X, train.y, test.X)
    loss[cv] $\leftarrow 1 - \frac{pred-test.y}{\text{length(pred)}}$
  end for
  avgLoss[k] $\leftarrow$ mean(loss)
end for
bestK $\leftarrow$ min(avgLoss)   ▷ k for which avgLoss is lowest.

The `shuffleSplit` method splits a dataset into n randomized disjoint subsets, in this case $n = 5$. The `bucketJoiner` method joins one or more of these disjoint subsets into a single, larger subset by appending one onto the other. The k-value with the lowest average loss for our data was 5, with an accuracy of approximately 80%. The accuracy on the test data is $68,4\%$.

1.3   I.4.3

The code for this experiment can be found in `I_4_3.m`. The experiment is very similar to that of `I.4.2`. Before cross-validating, we normalize the data using the method found in `scale.m`.

The μ values for the training data is $(5.7560, 0.3017)$ and $\sigma^2 = (0.6958, 0.0018)$. Normalizing the training data gives $\mu = (-0.3466 * 1^{-14}, -0.1823 * 1^{-14})$ and $\sigma^2 = (1.0000, 1.0000)$. After the normalization of the test data with the values found for the training data the test data have $\mu = (0.2073, 0.43)$ and $\sigma^2 = (1.0914, 1.2732)$. Running the test from `I.4.2` now gives $K_{\text{best}} = 1$ and an accuracy of $78, 95\%$ on the test data, and 100% accuracy on the training data. The last result is expected since when $K = 1$ any entry will be closest to itself (the euclidian distance is 0.)

References