Statistical Methods of Machine Learning
Assignment 3

Martin Grunbaum
Martin Jorgensen
Thomas Barnholdt

March 18, 2014

III.1.1

Derivative of $h(a)$

In the following we asume $a \neq 0$.

$$h(a) = \frac{a}{1 + |a|}$$
$$\left(\frac{f}{g}\right)' = \frac{f'g - g'f}{g^2}$$
$$f(a) = a$$
$$f'(a) = 1$$
$$g(a) = 1 + |a|$$
$$g^2(a) = (1 + |a|)^2$$
$$g'(a) = \frac{|a|}{a}$$
$$h'(a) = \frac{1(1 + |a|) - \frac{|a|}{a}a}{(1 + |a|)^2}$$
$$= \frac{1 + |a| - |a|}{(1 + |a|)^2}$$
$$= \frac{1}{(1 + |a|)^2}$$

III.1.2

What happens for small/large learning rates?

The smaller the learning rate, the less you adjust weights by and thus the smaller steps you take during each epoch of gradient descent. Conversely, the larger the learning rate, the bigger

a factor weights are adjusted by and the bigger steps you take during each epoch of gradient descent.

Beyond concerns about running time, a small learning rate may risk getting stuck in a very small local minimum and bigger learning rates risk bounding off past a good minimum. Selecting a good learning rate can thus be a case of trial and error, such as starting with a large learning rate to sketch out the general topology of the gradient and then switching to lower learning rates afterwards.


Figure 1: Mean-Squared-Error across epochs, for a Neural Network with 2 hidden nodes.


Figure 2: Mean-Squared-Error across epochs, for a Neural Network with 20 hidden nodes.

Figure 1 and Figure 2 show the Mean-Squared Error(MSE) of the two neural networks across

learning epochs. In both cases we see a steady decrease in error, however the error of the 20-node network is significantly lower than the 2-node network. It is probable that our 2-node network is incapable of illustrating the kind of distribution that the data stems from, causing severe issues for certain data points.

Sinc vs. Neural Network

Figures 3 and 4 show the `sinc(x)` function across the range $[-10; 10]$ alongside neural networks with 2 and 20 nodes respectively. The network with 2 nodes fits the data roughly like two functions of the kind $f(x) = a/(1 + norm(a))$ strung together with linear weight coefficients. Increasing the number of nodes to 20 allows the linear combination of weights and inputs and activations to express a distribution closer to that of `sinc(x)`.


Figure 3: sinc(x) vs. Neural Network with 2 hidden nodes, plotted across the range [-10; 10]

Overfitting and early stopping

Overfitting occurs when a fit becomes overly complex and tied to the data being used to train with. One solution to overfitting is to include a regularization parameter in order to lessen the effects of overfitting. Another solution is to perform early stopping: when training an optimization algorithm, the error measured with respect to independent data (such as a validation set) often shows a decrease at first, followed by an increase. We can stop at the point where the error on the validation set increases, in order to obtain a network that hopefully has better generalization performance than had we continued.

Figure 4: sinc(x) vs. Neural Network with 20 hidden nodes, plotted across the range [-10; 10]

III.2.1

Mean/variuance of the training data

priorTrainMean	priorTrainVar	trainMean	trainVar	testMean	testVar
	$1.0e+03*$	$1.0e-14*$			
155.9604	1.9830	0.0180	1.0000	-0.0782	0.7323
204.8212	9.7331	-0.0623	1.0000	-0.1572	0.7150
115.0586	2.1152	-0.0305	1.0000	0.0553	0.7977
0.0060	0.0000	0.0151	1.0000	0.1126	1.9906
0.0000	0.0000	-0.1268	1.0000	0.0712	1.6662
0.0032	0.0000	-0.0344	1.0000	0.0865	2.1370
0.0033	0.0000	0.0238	1.0000	0.1151	1.9225
0.0096	0.0000	-0.0019	1.0000	0.0866	2.1379
0.0277	0.0000	0.1217	1.0000	0.2477	1.7721
0.2624	0.0000	-0.0762	1.0000	0.2439	1.8292
0.0147	0.0000	0.1304	1.0000	0.2284	1.7175
0.0166	0.0000	-0.0491	1.0000	0.2496	1.7780
0.0220	0.0000	0.0597	1.0000	0.3150	2.1905
0.0440	0.0000	-0.0129	1.0000	0.2284	1.7176
0.0226	0.0000	0.0140	1.0000	0.1483	2.6633
22.0007	0.0167	-0.2299	1.0000	-0.0565	1.3610
0.4948	0.0000	0.1184	1.0000	0.0732	1.0827
0.7157	0.0000	0.3141	1.0000	0.0863	0.9514
-5.7637	0.0011	-0.1559	1.0000	0.1540	1.2166
0.2148	0.0000	0.0715	1.0000	0.3091	1.3629
2.3658	0.0001	-0.1414	1.0000	0.0870	1.1336
0.1997	0.0000	-0.0476	1.0000	0.1677	1.4149

III.2.2

Description of the software used:
We used LibSVM for Matlab[1]. Training the SVM produces a model structure, this is done like so:

```
model = svmtrain(trainX, trainY, 'kernel_function','rbf',
                 'rbf_sigma',sigmas,'boxconstraint',C);
```

trainX	Training data features.
trainY	Corresponding training data classes.
'kernel_function' & 'rbf'	Sets the kernel function to be a gaussian kernel function.
'rbf_sigma' & sigmas	Lets us specify a $\sigma = \sqrt{1/(2\gamma)}$.
'boxconstraint' & C	Lets us specify regularization paramter.

Our process:
For our five-fold cross validation we re-used the code (`shuffleSplit` and `bucketJoiner`) from the first assignment. So for each C and each γ we did a five-fold cross validation. The best C and Gamma combination was extracted and used to train a new model that could be used for classification.

Results:
We used the following numbers for C and γ: $C = (0.01, 0.1, 1.0, 10, 100, 1000, 10000)$ and $\gamma = (0.0001, 0.001, 0.01, 0.1, 1, 10, 100)$. For the raw data we found the best C and γ to be $C = 0.01, \gamma = 0.001$. For the normalized data we found $C = 100, \gamma = 0.1$.

Prenormalized training data accuracy	46.94%
Prenormalized test data accuracy	46.39%
Normalized training data accuracy	100.00%
Normalized test data accuracy	89.69%

Discussion:
If you do not normalize your data, you will unwillingly give more importance to feutures wich are widely spreed, and similarly give less importeanse to dense data.

III.2.3.1

For our optimum hyper-parameters the number of free support vectors are 0, and the number of bound support vectors are 54. Changing C will change the number of supportvectors in the model, a lower value of C will result in more supportvectors and with more free supportvectors. This means smaller C values gives you more supportvectors to work with, but the increase in free supportvectors also means that the machine is more prone to miss-classification.

---

[1] http://www.csie.ntu.edu.tw/~cjlin/libsvm/#matlab

III.2.3.2

As the number of entries in the training data increases so does the number of tsupportvectors in the SVM. The amount of free vectors (for our training dataset) is roughly one third of the total number of support vectors and so does not change with the amount of training data. The increase in the number of supportvectors will mean the SVM takes more time to process each query afterwards, but will also be able to do so with greater accuracy. So much like the C value, this becomes a trade-off, in this case between accuracy and speed.

References